

# FAF.PAD16.2 Autumn 2023

## Lab 2: Logically Linked DBs

**Handed out:** October 26, 2023

**Due:** December 13, 2023

### Checkpoints Terms

#### Checkpoint 1

Update your architecture with features you want to implement. Add in read.me a mini documentation about all your endpoints, how to run/deploy your project and steps to run/test docker images. Make clear what endpoint should be accessed first (if any), bodies/parameters for all of them. Also export Postman collection for all endpoints as json and upload it to Github.

#### Checkpoint 2

For Checkpoint 2, student is required to implement a part of the laboratory work, see Requirements section for tasks and marking system. Interpret it as being a Minimum Viable Product (MVP), and defend to the professor a functional or partially functional implementation. Professor is interested in student's progress!

The grading for this checkpoint will be specific. Students who miss the deadline will receive a grade of 1, which will be automatically recorded in the grade book. On the other hand, receiving a '+' means that the student is guaranteed a temporary mark. When student presents Checkpoint 3, this '+' will automatically convert into the grade received at the third checkpoint.

#### Checkpoint 3

Further details on the process of advance submission of the checkpoint and testing conditions will be provided by the professor a few days before the presentation day.

## Requirements

Mark	Team size: 1	Team size: 2
1	just be	(just be) x2
2	<ul style="list-style-type: none"> <li>• Mark 1</li> <li>• trip <b>Circuit Breaker</b> if multiple re-routes happen</li> </ul>	<ul style="list-style-type: none"> <li>• Mark 1</li> <li>• trip <b>Circuit Breaker</b> if multiple re-routes happen</li> </ul>
3	<ul style="list-style-type: none"> <li>• Mark 2</li> <li>• <b>Service High Availability</b></li> </ul>	<ul style="list-style-type: none"> <li>• Mark 2</li> <li>• <b>Service High Availability</b></li> </ul>
4	<ul style="list-style-type: none"> <li>• Mark 3</li> <li>• implement <b>ELK stack</b><a href="#">[3]</a> or <b>Prometheus</b><a href="#">[4]</a> + <b>Grafana</b><a href="#">[5]</a> for logging. Aggregate data from ALL services</li> </ul>	<ul style="list-style-type: none"> <li>• Mark 3</li> <li>• implement <b>ELK stack</b><a href="#">[3]</a> or <b>Prometheus</b><a href="#">[4]</a> + <b>Grafana</b><a href="#">[5]</a> for logging. Aggregate data from ALL</li> </ul>
5	<ul style="list-style-type: none"> <li>• Mark 4</li> <li>• implement microservice-based <b>2 Phase Commits</b> for a endpoint that create changes more than in one database (create new endpoint if needed).<a href="#">[1]</a></li> </ul>	<ul style="list-style-type: none"> <li>• Mark 4</li> <li>• implement microservice-based <b>2 Phase Commits</b> for a endpoint that create changes more than in two databases (create new endpoint if needed).<a href="#">[1]</a></li> </ul>
6	<ul style="list-style-type: none"> <li>• Mark 5</li> <li>• <b>Consistent Hashing for Cache</b> <a href="#">[7]</a></li> </ul>	<ul style="list-style-type: none"> <li>• Mark 5</li> <li>• <b>Consistent Hashing for Cache</b> <a href="#">[7]</a></li> </ul>
7	<ul style="list-style-type: none"> <li>• Mark 6</li> <li>• implement <b>Cache High Availability</b></li> </ul>	<ul style="list-style-type: none"> <li>• Mark 6</li> <li>• implement <b>Cache High Availability</b></li> </ul>
8	<ul style="list-style-type: none"> <li>• Mark 7</li> <li>• instead of <b>2 Phase Commits</b> implement <b>Long-running saga transactions</b> with coordinator <a href="#">[1]</a></li> </ul>	<ul style="list-style-type: none"> <li>• Mark 7</li> <li>• instead of <b>2 Phase Commits</b> implement <b>Long-running saga transactions</b> with coordinator <a href="#">[1]</a></li> </ul>
9	<ul style="list-style-type: none"> <li>• Mark 8</li> <li>• <b>Database redundancy/replication + failover</b> - implement any kind of replication for at least one database, minimum 4 replicas <a href="#">[2]</a></li> </ul>	<ul style="list-style-type: none"> <li>• Mark 8</li> <li>• <b>Database redundancy/replication + failover</b> - implement any kind of replication for at least two databases, minimum 4 replicas <a href="#">[2]</a></li> </ul>
10	<ul style="list-style-type: none"> <li>• Mark 9</li> <li>• create a <b>Data Warehouse</b> that will be periodically updated with all data from your databases. Use any ETL you want, it can be a job or separated service <a href="#">[9]</a></li> </ul>	<ul style="list-style-type: none"> <li>• Mark 9</li> <li>• create a <b>Data Warehouse</b> that will be periodically updated with all data from your databases. Implement your own ETL as separated service <a href="#">[9]</a></li> </ul>

## Readings

- [1] Keyang Xiang. “Patterns for distributed transactions within a microservices architecture”. <https://developers.redhat.com/blog/2018/10/01/patterns-for-distributed-transactions-within-a-microservices-architecture>.
- [2] Ben Lutkevich. “database replication”. <https://searchdatamanagement.techtarget.com/definition/database-replication>.
- [3] Elasticsearch. “What is the ELK Stack?”. <https://www.elastic.co/what-is/elk-stack>.
- [4] Cloud Native Computing Foundation. “Prometheus”. <https://prometheus.io>.
- [5] Grafana Labs. “Grafana”. <https://grafana.com>.
- [6] Redis. “Replication”. <https://redis.io/topics/replication>.
- [7] Juan Pablo Carzolio. “The Ultimate Guide to Consistent Hashing”. <https://www.toptal.com/big-data/consistent-hashing>.
- [8] l3a0. “Distributing a Cache”. <https://blog.baowebdev.com/2019/04/distributing-a-cache>.
- [9] “What is a data warehouse”. <https://www.ibm.com/topics/data-warehouse>.

**Good Luck!**